

CMT107 Visual Computing

VI.1 Texture Mapping


Xianfang Sun

School of Computer Science & Informatics
Cardiff University

Overview

- Texture mapping
 - Texture coordinates
 - Aliasing effects and MIP mapping
- Bump mapping
- Displacement mapping
- Light maps
- Shadow maps
- Texture Mapping in OpenGL

From Shading to Texture



Model



Model + Shading



Model + Shading
+ Textures


At what point
do things start
looking real?



For more info on the computer artwork of Jeremy Birn
see <http://www.3drender.com/jbirn/productions.html>

Texture

- Visual appearance of objects can be enhanced by textures
- The concept is simple



During rasterization interpolate the coordinate indices into the texture map

Texture Coordinates

- For each vertex specify *texture coordinates* $(s,t) \in [0,1]^2$
 - Canonical position of pixel in texture for vertex
 - For each point p on the 3D polygon, corresponding texture coordinates (s,t) are required
 - *Bilinearly interpolate* texture coordinates in 3D


- Texture coordinates for point on quad

$$p = sa + te$$

$$e = b + s(c - b)$$

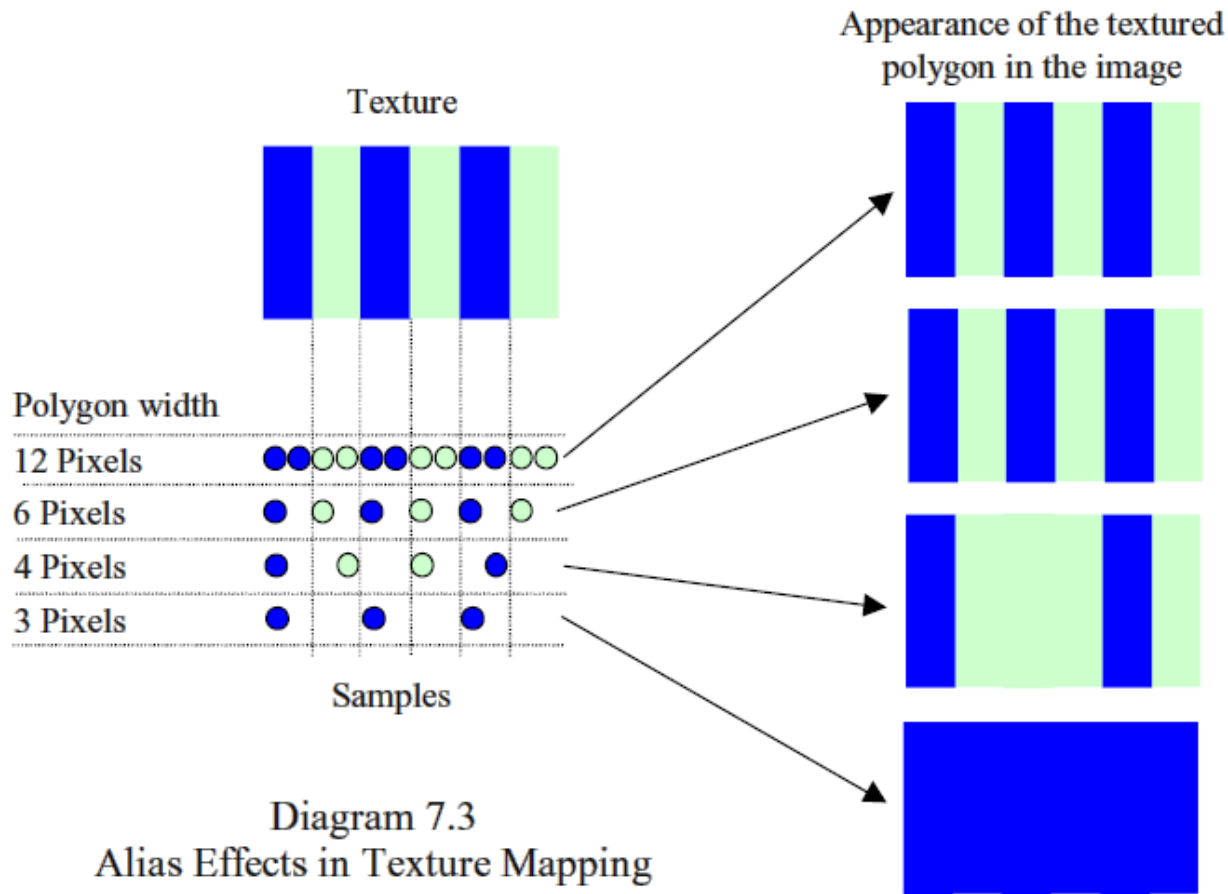
$$p = sa + tb + st(c - b)$$

→ Solve for (s,t) (assuming $(0,0)$ is texture coordinate of P_1)



Alias Effects

- One major problem of texture: *alias effects*
 - Caused by *undersampling*; results in *unreal artefacts*




Anti-aliasing

- Similar to untextured images use *anti-aliasing* technique
- Most successful approach: *supersampling*
 - Compute picture at a *higher resolution*
 - *Average* the supersamples to find pixel colour
 - This blurs boundaries, but leaves coherent areas of colour unchanged
 - Works well for polygons, but requires a lot of computations and does not work for line drawings
- Other approaches: convolution filtering
(see image processing)


MIP Mapping

- Popular technique of precomputing / prefiltering to address alias effects (MIP = multum in parvo; much in little)
- Basic idea: construct a *pyramid of images* for *different* texture sizes (prefiltered and resampled)
 - Pick texture image suitable for size or interpolate between texture images




Generalising Texture Mapping


- So far: texture is a *label* (colour) for each pixel
- Can use it to modify other things
 - E.g. use it for *illumination* to adjust material properties (all light types or only some of them)



Material



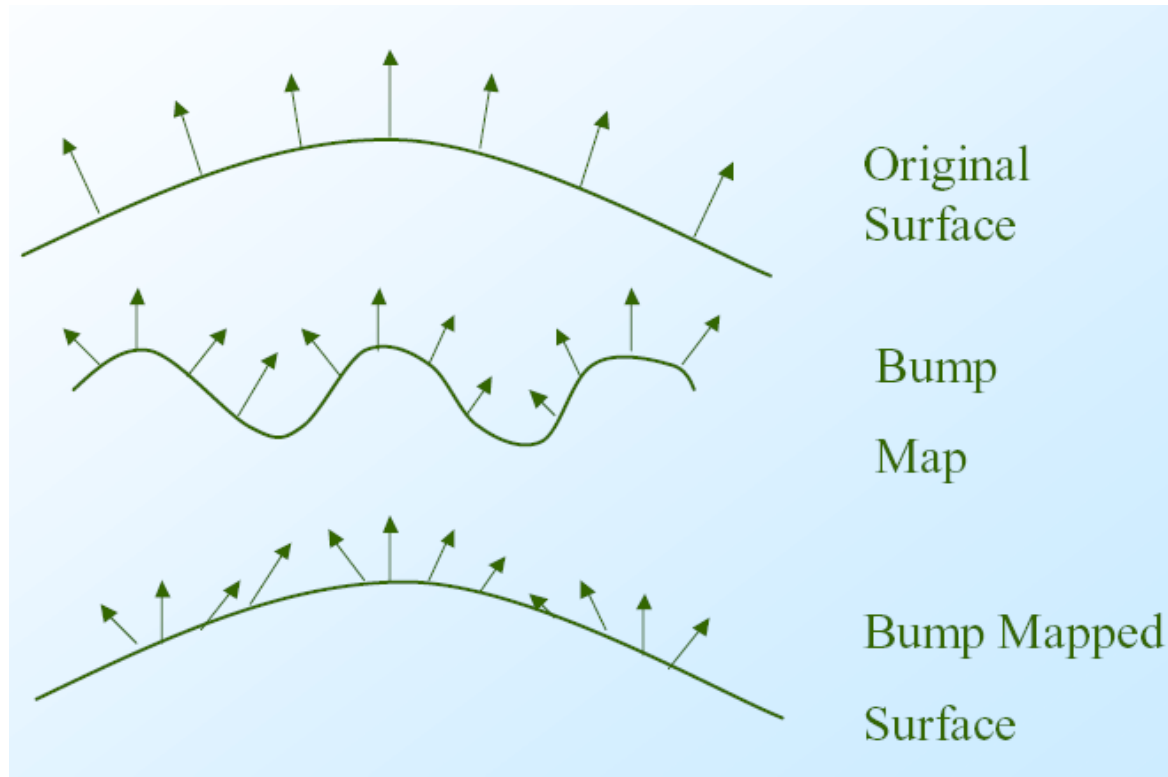
Texture as label




Texture as material

Bump Mapping


- Texture can be used to alter *surface normals* of an object
 - Does not change shape, but illumination computation
 - Changes in texture (partial derivatives) tell how to change the “height” of the normals



Bump Map Example



Sphere w/Texture





Bump Map



Bumpy Sphere


- As we do not change the shape, the silhouette does not change
 - Use only for small bumps
 - Requires illumination computation for each pixel (Phong shading, ray tracing, . . .)

Another Bump Map Example



Displacement Mapping

- Use texture map to *move* surface points



Light Maps

- In Quake *texture* and *light maps* are used
 - Light map contains precomputed illumination at low resolution
 - Multiply light map with texture map at run-time (and cache it)



Only Texture Map







Texture and Light Map

Shadow Maps

- Generate *shadows* using texture maps
 - Render scene from the *viewpoint of each light source* and only keep *depth buffer* values in *shadow buffers*
 - When shading each pixel (illumination computation per pixel):
 - Compute vector L from visible point to light source (needed for illumination computation)
 - Compute the length of L
 - *Compare* this length with corresponding value in the *shadow buffers*
 - If the shadow buffer value is *less*, then the point is in the shadow and we can ignore the light source

Shadow Map Example



Texture Mapping in OpenGL

- Use Texture and TextureIO to apply a texture
 1. Create a texture object using TextureIO
 - `TextureIO.newTexture(File, boolean);`
 2. Indicate how the texture is to be applied to each pixel
 - `Texture.setTextureParameteri (...)`
 3. Draw the scene, supplying both texture and geometric coordinates; send the coordinates to vertex shader, and send texture sampler to fragment shader
 - `Texture.getImageTexCoords().top() ...`

Texture Mapping in OpenGL

- Using OpenGL Core functions to apply a texture
 1. Create a texture object and specify a texture for that object
 - `glGenTextures(...)`
 - `glBindTexture(...)`
 - `glTexImage2D(...)`
 2. Indicate how the texture is to be applied to each pixel
 - `glTexParameteri(...)`
 3. Enable texture mapping
 - `glEnable(GL_TEXTURE_2D)`
 4. Draw the scene, supplying both texture and geometric coordinates; send the coordinates to vertex shader, and send texture sampler to fragment shader

- Step 0: Read in texture image

Texture Object

- Texture objects store texture data and keep it readily available for usage. Many texture objects can be generated.
- Generate identifiers for texture objects first
`int texids[n];`
`glGenTextures(n, texids)`
 - `n`: the number of texture objects identifiers to generate
 - `texids`: an array of unsigned integers to hold texture object identifiers
- Bind a texture object as the current texture
`glBindTexture(target, identifier)`
 - `target`: can be `GL_TEXTURE_1D`, `GL_TEXTURE_2D`, or `GL_TEXTURE_3D`
 - `identifier`: a texture object identifier
- Specify texture image
`glTexImage2D(target, level, internalFormat, width, height, border, format, type, data);`

Texture Object Example Code

```
int texids[] = new int[1];  
ByteBuffer texImg = readImage("textures/Day.png");  
  
glGenTextures(1, texids, 0);  
glBindTexture(GL_TEXTURE_2D, texids[0]);  
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGB,  
             texWidth, texHeight, 0, GL_BGR,  
             GL_UNSIGNED_BYTE, texImg);
```

Texture Parameters

- OpenGL has a variety of parameters that determine how texture is applied.
 - **Wrapping parameters** determine what happens if s and t are outside the $(0,1)$ range
 - **Filter modes** allow us to use area averaging instead of point samples
 - **Environment parameters** determine how texture mapping interacts with shading
 - **Mipmapping** allows us to use textures at multiple resolutions


- OpenGL Command

`glTexParameterf(target, pname, param);`


- **target**: Specifies the target texture
- **pname**: Specifies the symbolic name of a single-valued texture parameter
- **param**: Specifies the value of pname.

Wrapping Modes


- Repeat: use s, t modulo 1
- Clamp: if $s, t > 1$ use 1, if $s, t < 0$ use 0
 - `glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_REPEAT)`
 - `glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP_TO_EDGE)`
 - `GL_CLAMP_TO_BORDER, GL_MIRRORED_REPEAT...`



texture




GL_REPEAT
wrapping



GL_CLAMP_TO_EDGE
wrapping

Texture Filtering

- A pixel may be mapped to a small portion of a texel or a collection of texels from the texture map. How to determine the color of the pixel?
- **Magnification**: when a pixel mapped to a small portion of a texel
`glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, type);`
 - **type**: GL_NEAREST or GL_LINEAR
- **Minification**: when a pixel mapped to many texels
`glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, type);`
 - **type**: GL_NEAREST, GL_LINEAR, GL_NEAREST_MIPMAP_LINEAR, GL_LINEAR_MIPMAP_LINEAR, ...



Shading and Texture Interaction

- You can specify how the texture-map colors are used to modify the pixel colors by setting environment parameters in old version of OpenGL


```
glTexEnvf(GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE, mode);
```

mode values:

- **GL_REPLACE**: replace pixel color with texture color
- **GL_BLEND**: $C = C_f(1 - C_t) + C_c C_t$,
 - C_f is the pixel color, C_t is the texture color, and C_c is some constant color
- **GL_MODULATE**: $C = C_f C_t$ (Default)
- More on OpenGL programming guide
- In the shader version of OpenGL, the interaction should be implemented in the fragment shader.

Assign Texture Coordinates

- Every point on a surface should have a texture coordinate (s, t) in texture mapping
- We often specify texture coordinates to polygon vertices and interpolate texture coordinates with the polygon
- `Texture.getImageTexCoords()` can be used to retrieve texture coordinates



Typical Code in Main Program

```
// Set the texture to be used
try {
    texture = TextureIO.newTexture(new File("WelshDragon.jpg"), false);
} catch (IOException ex) {
    Logger.getLogger(getClass().getName()).log(Level.SEVERE, null, ex);
}

// Set texture coordinates
float[] texCoord = {...};
FloatBuffer textures = FloatBuffer.wrap(texCoord); gl.glGenBuffers(...);
gl.glBindBuffer(...);
gl.glBufferData(...);
gl.glBufferSubData(...);

// Send texture coordinates to vertex shader
vTexCoord = gl.glGetAttribLocation( program, "vTexCoord" );
gl.glEnableVertexAttribArray( vTexCoord );
gl.glVertexAttribPointer( vTexCoord, 2, GL_FLOAT, false, 0, offsetSize);

// Set the fragment shader texture sampler variable
gl.glUniform1i( gl.glGetUniformLocation(program, "tex"), 0 );
```

Vertex Shader

```
#version 330 core
```

```
layout(location = 0) in vec4 vPosition;  
layout(location = 1) in vec3 vColour;  
layout(location = 2) in vec2 vTexCoord;
```

```
out vec4 color;  
out vec2 texCoord;
```

```
uniform mat4 ModelView;  
uniform mat4 Projection;
```

```
void main()  
{  
    gl_Position = Projection * ModelView * vPosition;  
    texCoord    = vTexCoord;  
  
    color.rgb = vColour;  
    color.a = 1.0;  
}
```

Fragment Shader

```
#version 330 core

in  vec4 color;
in  vec2 texCoord;

out vec4 fColor;

uniform sampler2D tex;

void main()
{
    fColor = color* texture( tex, texCoord );
}
```

More details in the Labs...

Summary

- Describe the principle of texture maps. What are texture coordinates and how are they related to 3D coordinates?
- What options do exist to generalise texture maps? For what other effects are they useful and what are the advantages and disadvantages of these techniques?
- How to program texture mapping in OpenGL?

CMT107 Visual Computing

VI.2 Ray Tracing

Xianfang Sun

School of Computer Science & Informatics
Cardiff University

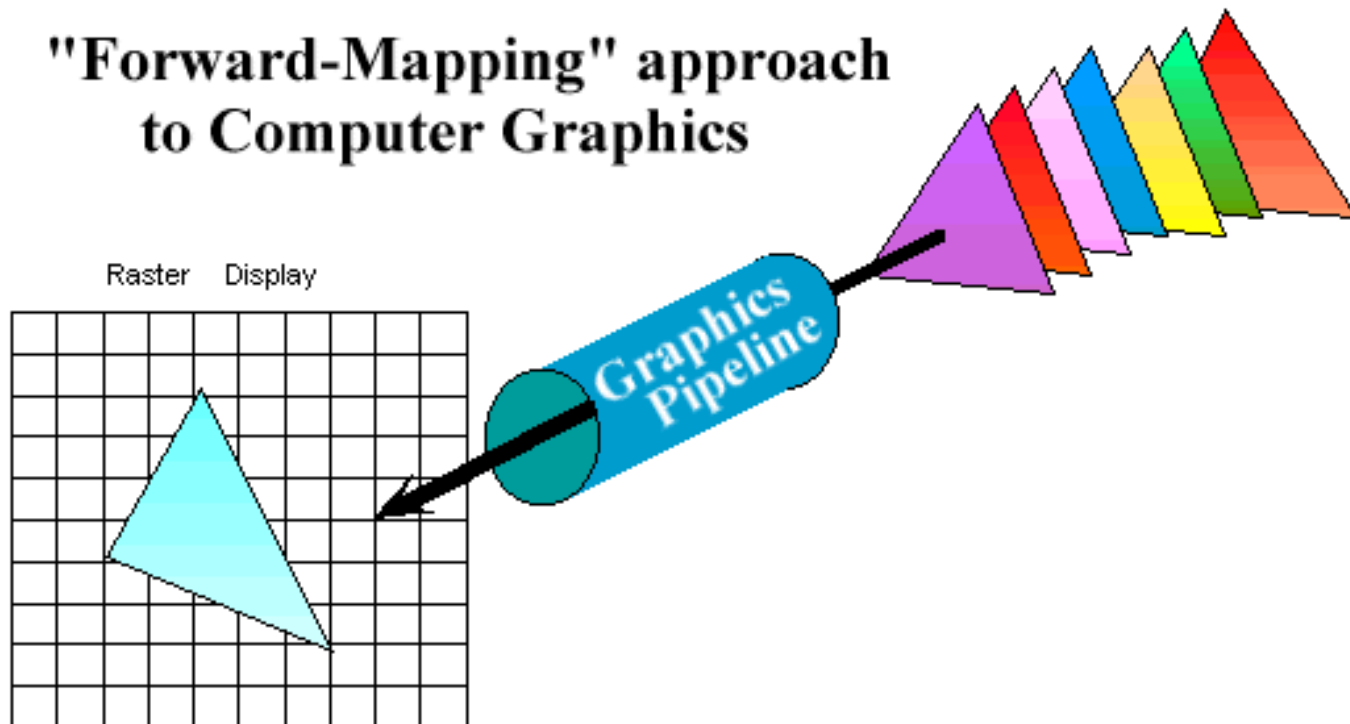
Overview

- Ray casting
- Ray tracing

Graphics Pipeline Review


- Properties of the graphics pipeline
 - Primitives are processed *one at a time* (in sequence)
 - All analytic processing done *early on*
 - Scan conversion (Rasterisation) occurs *last*
 - Minimal state required (*immediate mode* rendering)

**"Forward-Mapping" approach
to Computer Graphics**




Ray Casting

- An alternative to pipeline approach: *ray casting*
 - Search along lines of sight (*rays*) for visible primitive(s)
- Properties:
 - Go through *all primitives* at *each pixel*
(must have all primitives in a display list)
 - Sample (rasterisation) *first*
 - Do analytic processing *later*
- *Inverse mapping* approach



Global Illumination


- Ray casting properties:
 - Takes no advantage of *screen* space coherence
 - Requires *costly visibility computation*
 - Forces *per pixel illumination* evaluations
 - Not suited for immediate mode rendering
- In 1980 T. Whitted introduced *recursive ray casting (ray tracing)* to address global illumination



Ray Tracing


➤ For each ray from the viewing position:

- Compute *visible* object along the ray
- Compute *visibility* of each light source from the visible surface point using a new ray
- If there is an object between the surface point and the light source, *ignore* the light source; otherwise, Phong illumination model is used to evaluate the light intensity
- Can easily add reflection and refraction, etc.




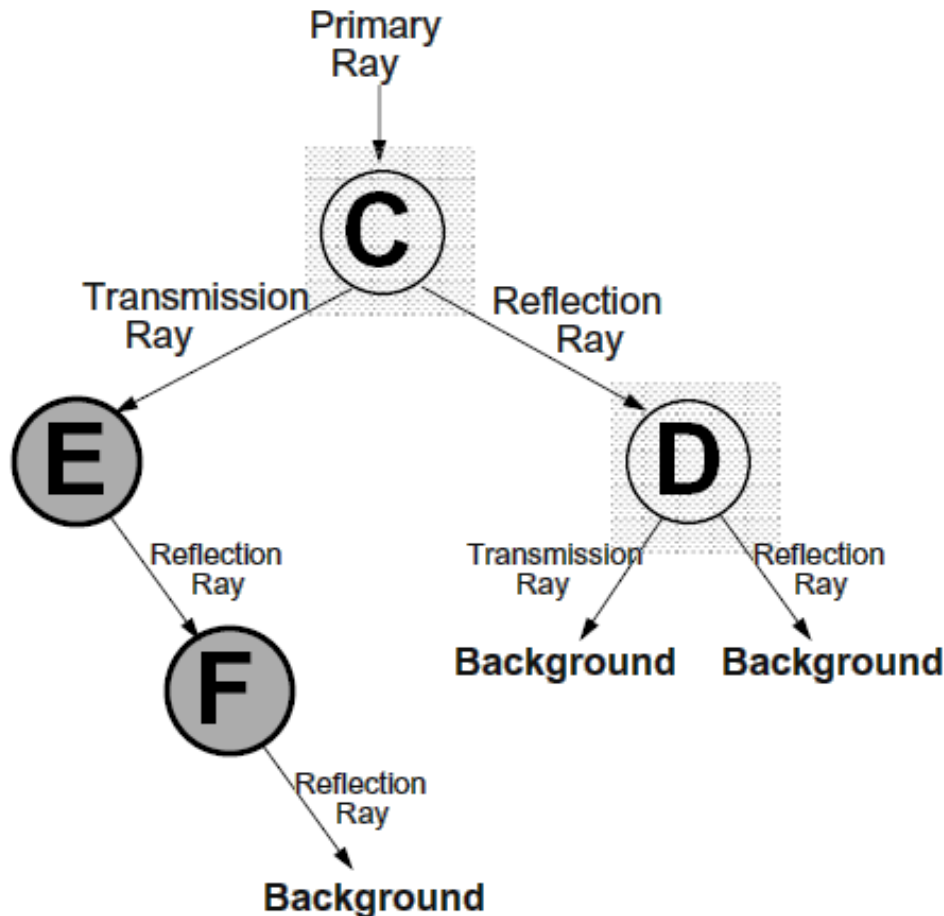
Ray Tracing

- For each object we need to know how to
- *reflect* light (Phong's illumination model)
 - *refract* light (Snell's law)
 - *emit* light (for light sources)
 - *intersect* object with a ray



Ray Tracing Tree

- Move up backwards in tree and **combine intensities** as determined by *Phong's illumination model*




From Pixels to Rays

- Compute **ray direction** $v(x, y)$ for raster coordinates (x, y)

- $u = \frac{\text{look} \times \text{up}}{\|\text{look} \times \text{up}\|}$

- $v = \frac{\text{look} \times u}{\|\text{look} \times u\|}$



- $o = \frac{\text{look}}{\|\text{look}\|} \frac{\text{width}}{2 \tan\left(\frac{\text{fov}}{2}\right)} - \frac{\text{width}}{2} u - \frac{\text{height}}{2} v$


- $v(x, y) = (xu_x + yv_x + o_x; xu_y + yv_y + o_y; xu_z + yv_z + o_z)$

Ray-Plane/Polygon Intersection

➤ *plane-line intersection*


- Ray: $P = P_0 + tV$
- Plane: $P^T N + D = 0$
- Substitute: $(P_0 + tV)^T N + D = 0$
- Solution: $t = - (P_0^T N + D) / (V^T N)$

➤ For intersection with polygon, check if intersection point lies inside polygon




Ray-Sphere Intersection

- *Intersect* a sphere with the ray (*algebraic*)
- Ray parameterisation: $P(t) = P_0 + tV$
 - Sphere equation: $\|P - O\|^2 - r^2 = 0$
 - Substitute: $\|P_0 + tV - O\|^2 - r^2 = 0$
 - Solve: $t^2 + 2V^T(P_0 - O)t + (\|P_0 - O\|^2 - r^2) = 0$



Ray-Sphere Intersection

- *Intersect* a sphere with the ray (*geometric*)
- $L = O - P_0$, $t_{ca} = L^T V$
 - if $t_{ca} < 0$, no intersection
 - $d^2 = L^T L - t_{ca}^2$
 - if $d > r$, no intersection
 - $t_{hc} = \sqrt{r^2 - d^2}$, $\rightarrow t = t_{ca} - t_{hc}$ and $t_{ca} + t_{hc}$




Ray Tracing Summary

- Input: viewing position \mathbf{v} , look-at point \mathbf{a} , up vector \mathbf{u}
- For each pixel:
 - Create a ray \mathbf{l} from the viewing position \mathbf{v} in direction \mathbf{d} such that it passes the pixel in the viewing plane
 - Set the colour to be the return value of $\text{raytrace}(\mathbf{v}, \mathbf{d})$
- Function $\text{raytrace}(\mathbf{v}, \mathbf{d})$:
 - Initialise position \mathbf{t} on ray \mathbf{l} from \mathbf{v} in direction \mathbf{d} to infinity and the nearest object \mathbf{n} to empty
 - For each object \mathbf{o} in the scene
 - Compute intersection \mathbf{p} of \mathbf{l} and \mathbf{o} closest to \mathbf{v}
 - If \mathbf{p} exists and is closer to \mathbf{v} than \mathbf{t} , set \mathbf{t} to \mathbf{p} and \mathbf{n} to \mathbf{o}
 - If \mathbf{n} is empty, return background colour, **else ...**

Ray Tracing Summary (cont.)

- **else ...**
 - If \mathbf{n} is reflective and we haven't reached the maximum recursion depth level, compute perfect reflection vector \mathbf{r} of \mathbf{d} at \mathbf{t} and call $\text{raytrace}(\mathbf{t}, \mathbf{r})$ to obtain reflected colour c_r
 - If \mathbf{n} is transparent and we haven't reached the maximum recursion depth level, compute refraction vector \mathbf{r}' of \mathbf{d} at \mathbf{t} and call $\text{raytrace}(\mathbf{t}, \mathbf{r}')$ to obtain refracted colour c_t
 - For each light source $k=1, \dots, m$ at position \mathbf{l}_k , cast ray from \mathbf{t} to \mathbf{l}_k . If this line segment intersects with any of the other objects, \mathbf{t} is in the shadow of this object. Otherwise compute the amount of light c_k reaching \mathbf{t} from k
 - Return combination of colours c_r , c_t and c_k , $k=1, \dots, m$

Examples



Properties of Ray Tracing

➤ *Advantages*

- Improved realism (shadows, reflections, transparency)
- Higher level rendering primitives
- Very simple design

➤ *Disadvantages*



- Very slow per pixel calculations
- Only approximate global illumination
(cannot follow all rays)
- Hard to accelerate with hardware

➤ *Acceleration* approach

- Try to reduce number of intersection computations


Ray Tracing Acceleration

- *Bounding volumes*
 - Check simple bounding volume for ray/surface intersections before checking complex shapes
- *Bounding volume hierarchies*
 - Construct and check hierarchical bounding volumes




Spatial Data Structures

- Create a data structure aware of the spatial relations
 - *Partition* space and place objects within subregions
 - *Only* consider subregions that the ray passes through
 - *Avoid computing intersections twice* if object lies inside multiple subregions




BSP Trees in Ray Tracing

- Partition space using *BSP Tree*




Rendered Examples



Advanced Phenomena


- Ray tracers can simulate (not always efficiently)
 - Soft shadows
 - Fog
 - Frequency dependent light
(Snell's law is different for different wave-lengths)
- But can barely handle diffuse/ambient lighting
 - *Radiosity* is a global illumination scheme complementing ray-tracing for diffuse/ambient lighting

Examples



(by Oliver Kreylos, <http://graphics.cs.ucdavis.edu/~okreylos/Private/RaytracingCorner/>)

Examples




(by Oliver Kreylos, <http://graphics.cs.ucdavis.edu/~okreylos/Private/RaytracingCorner/>)

Examples



(by Oliver Kreylos, <http://graphics.cs.ucdavis.edu/~okreylos/Private/RaytracingCorner/>)

Examples



(by Oliver Kreylos, <http://graphics.cs.ucdavis.edu/~okreylos/Private/RaytracingCorner/>)

Summary

- What is ray casting? What are its advantages and disadvantages?
- What is ray tracing? What are its advantages and disadvantages?
- How can we compute the rays through raster points for ray tracing? How can we compute the intersections of such a ray and a plane or a sphere? How is this done for other shapes?
- How can ray tracing be accelerated?