

CMT107 Visual Computing

VII.1 Image Processing in Java

Xianfang Sun

School of Computer Science & Informatics
Cardiff University

Overview

- Images in Java
- Read (Load) an Image
- Draw an Image
- Process an Image
- Write (Save) an Image

Images in Java

- An image is typically a rectangular two-dimensional array of pixels
 - each pixel represents the colour at that position
 - dimensions represent the horizontal extent (width) and vertical extent (height) of the image as it is displayed
- Most important image class in Java 2D API
 - `java.awt.image.BufferedImage`
- Image programming Tasks:
 - Load an external image file
 - Draw an image onto a drawing surface
 - Manipulate the pixels of an image
 - Save the contents of an image to an external image file

Read (Load) an Image

- Use javax.imageio package

```
BufferedImage img = null;
```

```
try {
```

```
    img = ImageIO.read(new File("Daffodil.jpg"));
```

```
} catch (IOException e) {
```

```
}
```

Draw An Image

➤ Use Graphics Function

```
boolean Graphics.drawImage(Image img, int x, int y,  
                             ImageObserver observer);
```

➤ Example

```
public void paint(Graphics g) {  
    g.drawImage(img, 0, 0, null);  
}
```

Process an Image

- The width and height of the image can be obtained by

```
width = img.getWidth();  
height = img.getHeight();
```

- The pixel colour at (x, y) can be retrieved and set by

```
Color pixel = new Color(img.getRGB(x, y));  
img.setRGB(x, y, pixel.getRGB());
```

- Example: convert a colour image to a grayscale image

```
for (int y = 0; y < height; y++)  
    for (int x = 0; x < width; x++) {  
        Color pixel = new Color(in.getRGB(x, y));  
        int r = pixel.getRed();  
        int g = pixel.getGreen();  
        int b = pixel.getBlue();  
        r = g = b = (int) (0.299*r + 0.587*g + 0.114*b); //grayscale  
        out.setRGB(x, y, (new Color(r, g, b)).getRGB());  
    }  
}
```

Write (Save) an Image

- Use javax.imageio package

```
BufferedImage out = getMyImage(); //Retrieve image  
try {
```

```
    ImageIO.write(out, "jpg", new File("DaffodilG.jpg"));  
} catch (IOException ex) {  
}
```

Summary

- What is an image?
- What is a pixel?
- How to load and save an image?
- How to draw an image?
- How to access and set the pixels of an image?

CMT107 Visual Computing

VII.2 Image Filtering

Xianfang Sun

School of Computer Science & Informatics
Cardiff University

Overview

- Linear filtering
- Convolution
- Box Filtering
- Gaussian Filtering
- Separable Kernel
- Median Filter
- Sharpening


Acknowledgement

The majority of the slides in this section are from Svetlana Lazebnik at University of Illinois at Urbana-Champaign

Image Filtering

- **Filtering** is a technique for **modifying** or **enhancing** an image.
 - Emphasise certain features or remove other features
- Filtering is a **neighbourhood operation**
 - The output value of any given pixel is determined by the values of the pixels in the neighbourhood of the corresponding input pixel
- **Linear filtering** is filtering in which the value of an output pixel is a **linear combination (weighted average)** of the values of the pixels in the input pixel's neighbourhood
 - Linear filtering can be represented by convolution

Linear Filtering



Motivation: Image Denoising

- How can we reduce noise in a photograph?



Moving average

- Let's replace each pixel with a *weighted* average of its neighborhood
- The weights are called the *filter kernel*
- What are the weights for the average of a 3x3 neighbourhood?

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1


“box filter”

Convolution

- Let f be the image and g be the kernel. The output of convolving f with g is denoted $f * g$.

$$(f * g)[m, n] = \sum_{k, l} f[m - k, n - l] g[k, l]$$

Convention:
kernel is “flipped”



Key properties




- **Linearity:** $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$
- **Shift invariance:** same behavior regardless of pixel location: $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$
- Theoretical result: any linear shift-invariant operator can be represented as a convolution

More Properties

- **Commutative:** $a * b = b * a$
 - Conceptually no difference between filter and signal
- **Associative:** $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- **Distributive** over addition: $a * (b + c) = (a * b) + (a * c)$
- **Scalars factor out:** $ka * b = a * kb = k(a * b)$
- **Identity:** unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$, $a * e = a$

Size of the Output


- 'full': output size is the sum of sizes of f and g - 1
- 'same': output size is the same as f
- 'valid': output size is the difference of the sizes of f and g



Boundary Pixels

➤ What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Practice with linear filters



Original

0	0	0
0	1	0
0	0	0


?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



Original

0	0	0
1	0	0
0	0	0

?

Practice with linear filters



Original

0	0	0
1	0	0
0	0	0



Shifted *left*
By 1 pixel

Practice with linear filters



Original

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1


?

Practice with linear filters



Original

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Blur (with a
box filter)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

−

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1





Sharpening filter

- Accentuates differences with local average


Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?



Smoothing with box filter revisited



- What's wrong with this picture?
- What's the solution?
 - To eliminate edge effects, weight contribution of neighbourhood pixels according to their closeness to the center



“fuzzy blob”

Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5, $\sigma = 1$


- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$ with 30×30
kernel





$\sigma = 5$ with 30×30
kernel

- Standard deviation σ : determines extent of smoothing


Choosing kernel width

- The Gaussian function has infinite support, but discrete filters use finite kernels




Choosing kernel width

- Rule of thumb: set filter half-width to about 3σ



Gaussian vs. box filtering



Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convolving two times with Gaussian kernel with std. dev. σ is same as convolving once with kernel with std. dev. $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:

Why is separability useful?

- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
- What if the kernel is separable?
 - $O(n^2 m)$

Noise



Original



Salt and pepper noise



Impulse noise




Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

Gaussian Noise

- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise

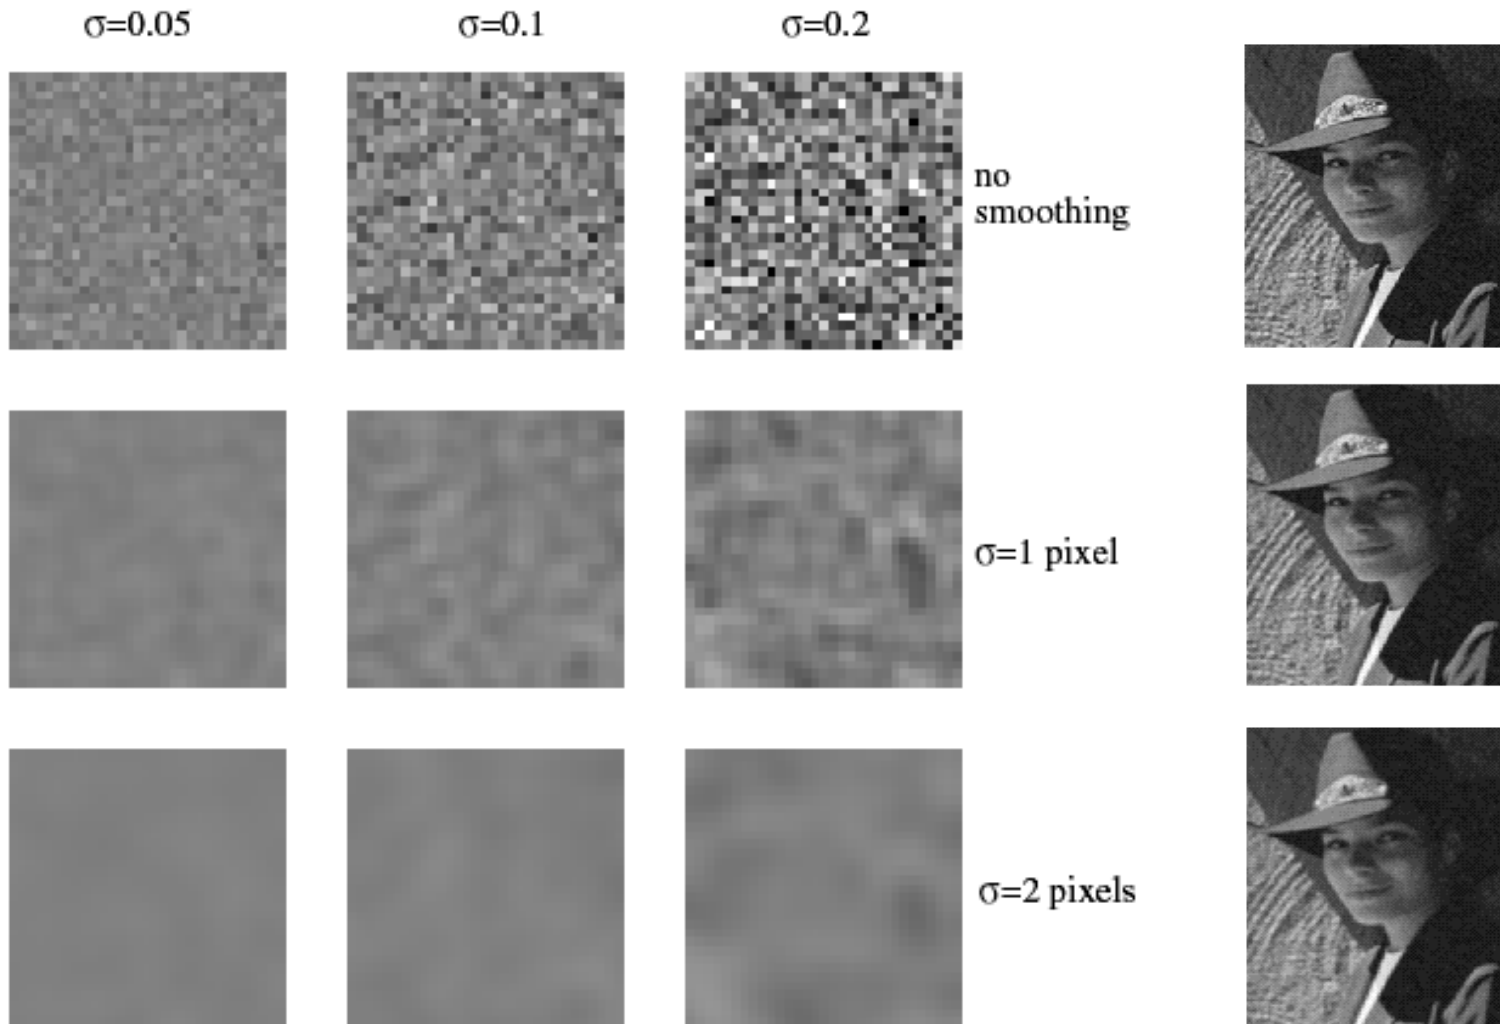
Image
Noise



$$f(x, y) = \underbrace{\hat{f}(x, y)}_{\text{Ideal Image}} + \underbrace{\eta(x, y)}_{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Reducing salt-and-pepper noise

3x3



5x5




7x7



➤ What's wrong with the results?

Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window




- Is median filtering linear?

Median filter

➤ What advantage does median filtering have over Gaussian filtering?


- Robustness to outliers

filters have width 5 :





Median filter

Salt-and-pepper noise



Median filtered



Gaussian vs. median filtering

Gaussian

3x3



5x5




7x7




Median



Sharpening revisited



before



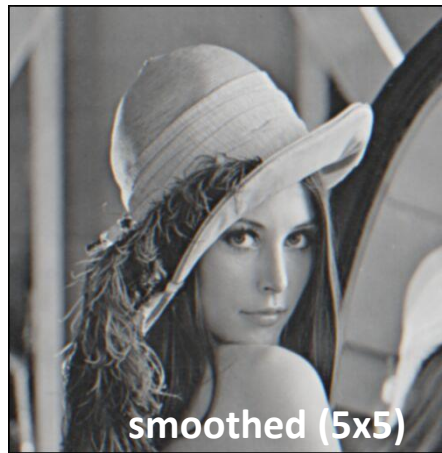
after

Sharpening revisited


➤ What does blurring take away?



−




=



Let's add it back:



+ α



=



Unsharp mask filter

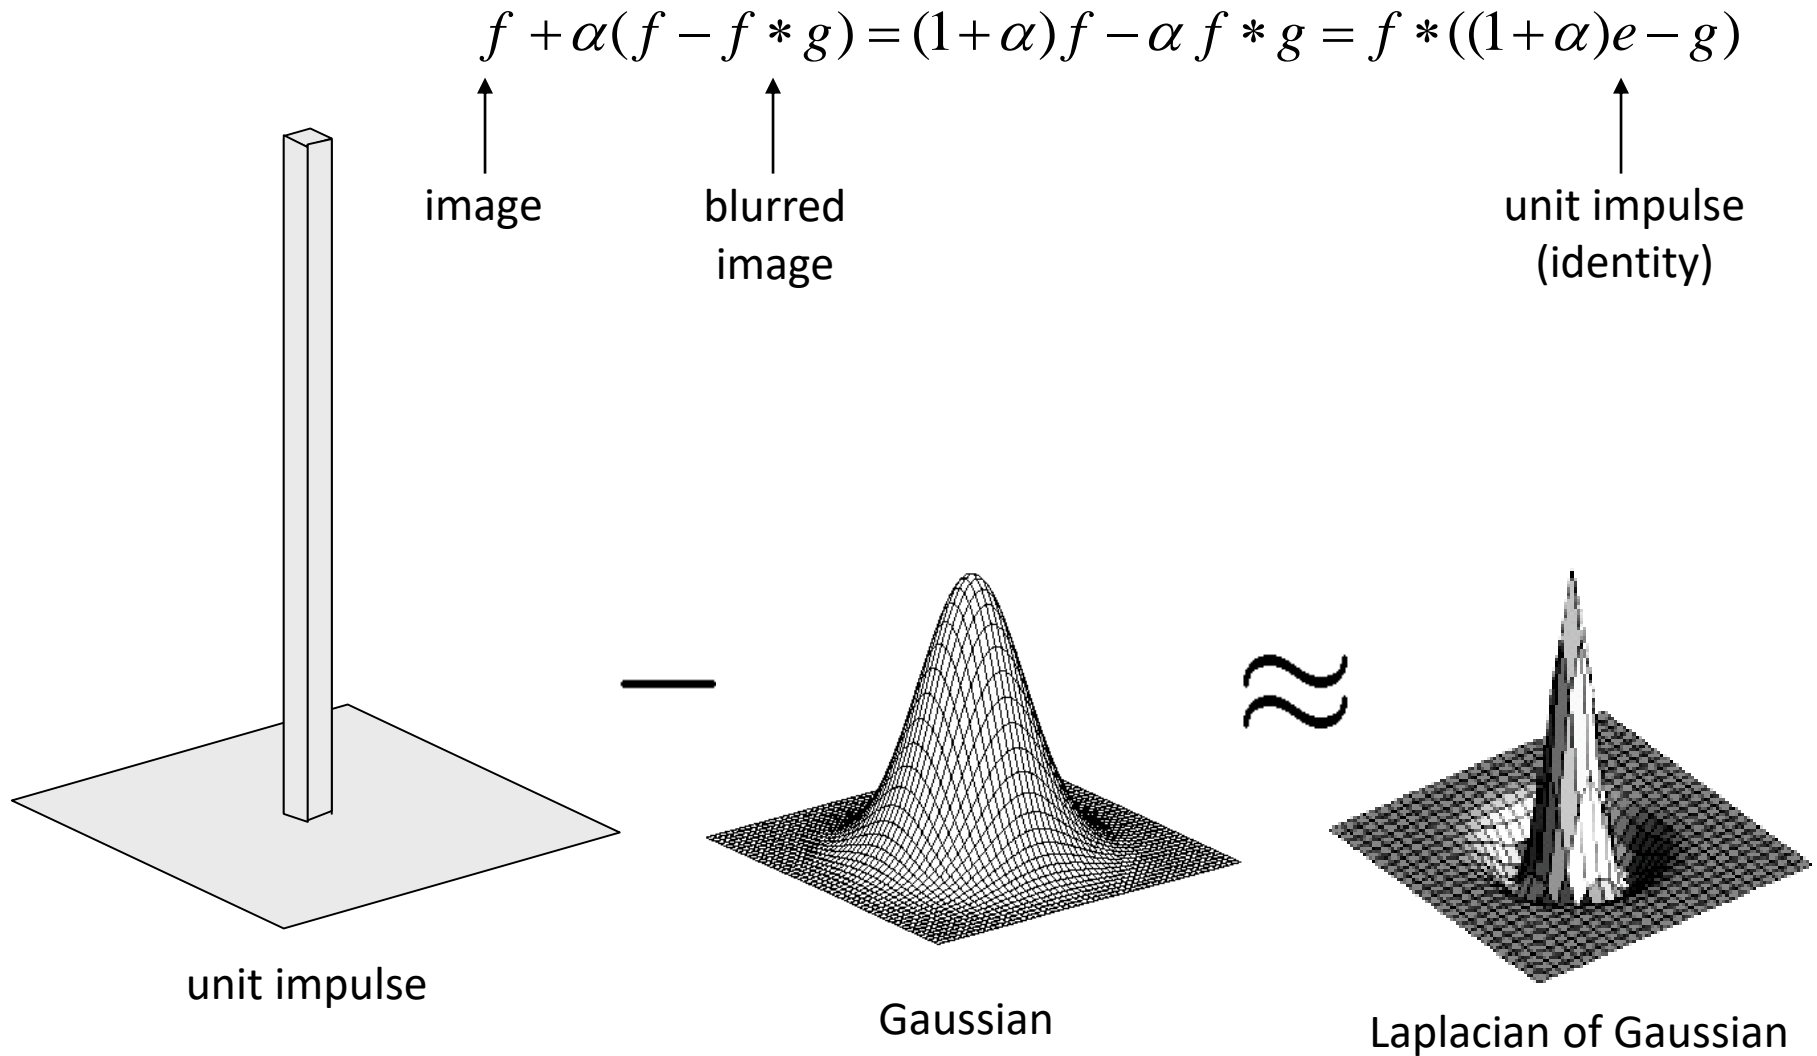


Image Filtering with Java

- Use filter() function in BufferedImageOp
- Implement filtering without using filter() function

Use filter() Function

➤ Define a filter kernel

```
float[] km = {          // low-pass filter kernel
    0.1f, 0.1f, 0.1f,
    0.1f, 0.2f, 0.1f,
    0.1f, 0.1f, 0.1f
};
Kernel kernel = new Kernel(3, 3, km);
```

➤ Define an Operator

```
BufferedImageOp op = null;
op = new ConvolveOp(kernel, ConvolveOp.EDGE_NO_OP, null);
```

- ConvolveOp(Kernel kernel, int edgeCondition, RenderingHints hints)
— edgeCondition: ConvolveOp.EDGE_NO_OP or
ConvolveOp.EDGE_ZERO_FILL

➤ Call the filter() function

```
out = new BufferedImage(width, height, BufferedImage.TYPE_INT_RGB);
op.filter(in, out);
```

Not Use filter() Function

➤ Define a filter kernel matrix

```
float[] km = {           // low-pass filter kernel
    0.1f, 0.1f, 0.1f,    // Suppose the matrix has been flipped
    0.1f, 0.2f, 0.1f,
    0.1f, 0.1f, 0.1f
};
```

➤ Calculate convolution on each pixel

```
int[] rArray = new int[width*height]; //
for each pixel {
    get the neighbourhood colours of the pixel
    calculate the colour according to the convolution formula
    set the pixel colour in the output image
}
```

- More details in Lab session 6

Summary

- What is filtering? What is linear filtering?
- What is convolution?
- How to do sharpening of image?
- What is box filtering, Gaussian filtering, and median filtering?
- What is separable kernel? Why use separable kernel?

CMT107 Visual Computing

VII.3 Corner Extraction

Xianfang Sun

School of Computer Science & Informatics
Cardiff University

Overview

- Feature Extraction
 - Characteristics of Good Features
 - Applications
- Corner Detection
 - Basic Idea
 - Mathematics
- Harris Detector
- Invariance and Covariance

Acknowledgement

The majority of the slides in this section are from Svetlana Lazebnik at University of Illinois at Urbana-Champaign

Feature Extraction: Corners



Why Extract Features?

➤ Motivation: panorama stitching


- We have two images – how do we combine them?



Why Extract Features?

➤ Motivation: panorama stitching

- We have two images – how do we combine them?



Step 1: extract features

Step 2: match features

Why Extract Features?

➤ Motivation: panorama stitching

- We have two images – how do we combine them?




Step 1: extract features

Step 2: match features

Step 3: align images

Characteristics of Good Features



➤ Repeatability

- The same feature can be found in several images despite geometric and photometric transformations

➤ Saliency

- Each feature is distinctive

➤ Compactness and efficiency

- Many fewer features than image pixels

➤ Locality

- A feature occupies a relatively small area of the image; robust to clutter and occlusion

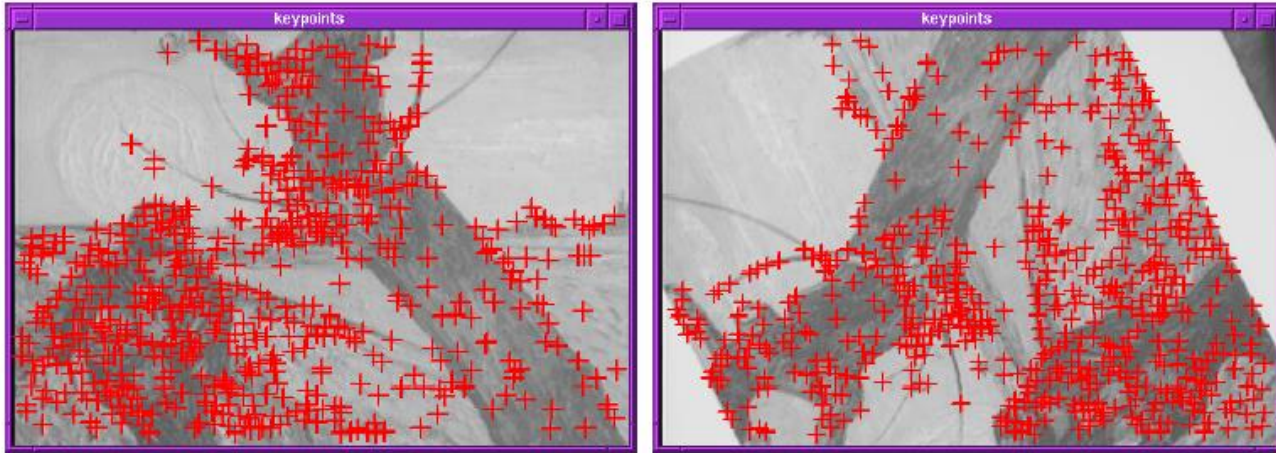
Applications

➤ Feature points are used for:

- Image alignment
- 3D reconstruction
- Motion tracking
- Robot navigation
- Indexing and database retrieval
- Object recognition



Finding Corners




- Key property: in the region around a corner, image gradient has two or more dominant directions
- Corners are repeatable and distinctive


C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)
Proceedings of the 4th Alvey Vision Conference, 1988: pages 147--151.

Corner Detection: Basic Idea


- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity



“flat” region:
no change in
all directions



“edge”:
no change
along the edge
direction




“corner”:
significant
change in all
directions

Corner Detection: Mathematics


Change in appearance of window $w(x,y)$
for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

$I(x, y)$





$E(u, v)$



Corner Detection: Mathematics

Change in appearance of window $w(x,y)$
for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$



Corner Detection: Mathematics

Change in appearance of window $w(x,y)$
for the shift $[u,v]$:


$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Window
function


Shifted
intensity

Intensity

Window function $w(x,y) =$



1 in window, 0 outside



or

Gaussian


Corner Detection: Mathematics

Change in appearance of window $w(x,y)$
for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$E(u, v)$



Corner Detection: Mathematics

Change in appearance of window $w(x,y)$
for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

Local quadratic approximation of $E(u,v)$ in the neighborhood of $(0,0)$ is given by the *second-order Taylor expansion*:

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

Corner Detection: Mathematics

$$E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about $(0, 0)$:

$$E(u, v) \approx E(0, 0) + [u \ v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{uv}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_u(u, v) = \sum_{x, y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_x(x+u, y+v)$$

$$E_{uu}(u, v) = \sum_{x, y} 2w(x, y) I_x(x+u, y+v) I_x(x+u, y+v) \\ + \sum_{x, y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xx}(x+u, y+v)$$

$$E_{uv}(u, v) = \sum_{x, y} 2w(x, y) I_y(x+u, y+v) I_x(x+u, y+v) \\ + \sum_{x, y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xy}(x+u, y+v)$$

Corner Detection: Mathematics

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about $(0, 0)$:

$$E(u, v) \approx E(0, 0) + [u \ v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{uv}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0, 0) = 0$$

$$E_u(0, 0) = 0$$

$$E_v(0, 0) = 0$$

$$E_{uu}(0, 0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0, 0) = \sum_{x, y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0, 0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_y(x, y)$$

Corner Detection: Mathematics

$$E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about $(0, 0)$:

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} \sum_{x, y} w(x, y) I_x^2(x, y) & \sum_{x, y} w(x, y) I_x(x, y) I_y(x, y) \\ \sum_{x, y} w(x, y) I_x(x, y) I_y(x, y) & \sum_{x, y} w(x, y) I_y^2(x, y) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0, 0) = 0$$

$$E_u(0, 0) = 0$$

$$E_v(0, 0) = 0$$

$$E_{uu}(0, 0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0, 0) = \sum_{x, y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0, 0) = \sum_{x, y} 2w(x, y) I_x(x, y) I_y(x, y)$$

Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a *second moment matrix* computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$


$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

Interpreting the Second Moment Matrix

The surface $E(u,v)$ is locally approximated by a quadratic form. Let's try to understand its shape.

$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



Interpreting the Second Moment Matrix

First, consider the axis-aligned case (gradients are either horizontal or vertical)


$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

If either λ is close to 0, then this is **not** a corner, so look for locations where both are large.

Interpreting the Second Moment Matrix

Consider a horizontal “slice” of $E(u, v)$:
$$\begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

This is the equation of an ellipse.




Interpreting the Second Moment Matrix

Consider a horizontal “slice” of $E(u, v)$:
$$\begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$


This is the equation of an ellipse.

Diagonalization of M :
$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$


The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by R



Visualization of Second Moment Matrices




Visualization of Second Moment Matrices



Interpreting the Eigenvalues


Classification of image points using eigenvalues of M :



Corner response function

$$R = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

α : constant (0.04 to 0.06)




Harris Detector: Steps

1. Compute Gaussian derivatives at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel
3. Compute corner response function R
4. Threshold R
5. Find local maxima of response function (nonmaximum suppression)


C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#) *Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

Harris Detector: Steps




Harris Detector: Steps

Compute corner response R



Harris Detector: Steps

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Steps

Take only the points of local maxima of R



Harris Detector: Steps




Invariance and Covariance

- We want corner locations to be *invariant* to photometric transformations and *covariant* to geometric transformations
- **Invariance:** image is transformed and corner locations do not change
 - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations





Affine Intensity Change




$$I \rightarrow aI + b$$

- Only derivatives are used => invariance to intensity shift $I \rightarrow I + b$
- Intensity scaling: $I \rightarrow aI$



Partially invariant to affine intensity change


Image Translation



- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation


Image Rotation



- Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

Scaling



Corner location is not covariant to scaling!

Summary

- Why we need feature extraction? What are the applications of feature extraction?
- What are Characteristics of Good Features?
- Describe the basic idea of corner detection.
- How to decide whether a point is in a flat region, on an edge, or corner according to the two eigenvalues of the second moment matrix?
- Describe steps of Harris detector
- What is Invariance and Covariance?
- Is affine intensity change invariant? Is image translation, rotation, or scaling covariant?